

Real-Time Object Detection System Report

<https://github.com/manasjuneja/objectrecognition>

Implementation Approach:

- **Model Selection:**
 - Used Ultralytics YOLOv8n (nano) pre-trained model for optimal speed and accuracy.
- **Video Processing:**
 - Captured real-time video frames from webcam using OpenCV.
- **Detection Pipeline:**
 - Performed inference on each frame; drew bounding boxes, class labels, and confidence scores.
- **Cross-Platform Compatibility:**
 - Ensured code runs on Windows, Linux, macOS, x86_64, and ARM devices.
- **Performance Optimization:**
 - Reduced input resolution and used model fusion for faster inference.

Results & Performance Analysis:

- **Real-Time FPS:**
 - Achieved 25-35 FPS on modern CPUs; up to 50+ FPS with GPU acceleration.
- **Detection Quality:**

- Detected 60+ COCO classes with high accuracy in real-time.
- **Accuracy Metrics:**
 - Evaluated on COCO/Roboflow datasets: $\text{mAP@0.5} \approx 37\text{--}45\%$ for YOLOv8n.
- **Resource Usage:**
 - Low memory and CPU usage; suitable for edge devices.

Challenges Faced & Solutions:

- **Low FPS on Some Machines:**
 - Low FPS (~4) on CPU initially - solved by switching to `yolov8n` for faster inference
 - Lowered webcam and model input resolution to 320x240.
 - Used GPU if available; optimized code for minimal overhead.
- **Accuracy Evaluation:**
 - Used public datasets with ground truth labels for proper metrics.
- **Cross-Platform Issues:**
 - Avoided OS-specific dependencies; tested on multiple platforms.

Future Improvements:

- **User Interface:**

- Add a simple GUI for toggling detection/class filters and viewing metrics.
- **Custom Training:**
 - Fine-tune YOLOv8 on domain-specific datasets for improved accuracy.
- **Deployment:**
 - Package as a Docker container or mobile app for broader accessibility.